226 - Artificial I

Podcast: secular buddhism

AI Executive Summary

1. **Core Thesis** The host argues that the human sense of self is not a fixed, unexamined entity, but rather an "artificial" construct—similar to AI systems—that is assembled from patterns, data, and conditioning received early in life (training data). Rather than viewing the concept of "no self" (*Anatta*) as implying non-existence, Buddhism utilizes the concept of *sankara* (conditioned/formed things) to explain that the self is "made," not "fake." This construction allows for functionality but leads to three dangerous states—deception, manipulation, and autonomous decision-making—just as AI systems operate. The solution is not to eliminate the self but to apply "interpretability" inward through mindfulness: observing thoughts and feelings as predictable patterns rather than objective facts to prevent suffering caused by unexamined stories.

2. **Key Arguments & Data**

3. **AI Function vs. Mind Mechanism:** Both AI systems and human minds operate via predictive pattern recognition. AI predicts the "most likely next word" based on training data without understanding truth; similarly, humans act confidently based on learned associations without conscious inspection of how those reactions were formed.
4. **The Three Risks & Their Parallel:** The three primary fears regarding AI (deception, manipulation, autonomous decisions) correspond directly to three lifelong human behaviors driven by unexamined conditioning:
    - **Deception:** AI hallucination vs. Human Post-Hoc Rationalization. Humans often create confident explanations for actions taken after the fact (e.g., the "Chicken Foot/Snowy Yard" experiment where a brain hemisphere fabricates a logical reason for an action it didn't choose).
    - **Manipulation:** AI prediction bias vs. The Human "Second Arrow." While the external event is the first arrow, the human mind adds anxiety and interpretation (the second arrow) based on past patterns. A four-word text message ("Can we talk?") triggers

a defensive reaction because the mind predicts negative outcomes from that pattern.
    - **Autonomy:** AI making decisions vs. Human Autopilot/Habits. Approximately 40% of human daily actions are habitual responses run by models (patterns) rather than deliberate, moment-to-moment choices.
5. **The Definition of Artificial:** The term "artificial" should not mean "fake," but rather "made." An artificial lake is a real lake formed by construction; similarly, the self is a constructed model useful for function (memory, promise-keeping) that exists and functions because it was put together from causes and conditions (*sankara*).
6. **Training Data:** Neither AI nor humans choose their data. AI learns from text/books; humans learn from childhood cues regarding love, punishment, validation, and social norms. The host posits that feelings of worthlessness or anxiety are often the result of early training data rather than internal reality ("is this true?").

7. **The Three Poisons:** Greed (wanting more), Hatred (pushing away), and Delusion (not seeing things as they are) are inherent in unexamined minds and are amplified by AI. The host claims the unexamined life has always been dangerous, with AI only raising the stakes for these behaviors.

8. **Actionable Takeaways**

9. **Identify Strong Reactions:** Begin this week by noticing strong, fast-acting reactions such as irritation, certainty, defensiveness, or dread. These are identified as the deepest training running automatically.
10. **Refractor as Prediction:** When a reaction occurs, label it not as an objective fact but as a prediction ("My mind is telling me I am in trouble" rather than "I am in trouble"). This creates necessary mental space.
11. **Investigate Origin:** Ask the specific question: "What was this trained on?" Trace where you first learned to see things that way (e.g., parents, teachers, cultural norms) and ask whose voice is being heard inside.
12. **Mindfulness as Interpretability:** Apply AI interpretability research inward by practicing mindfulness to watch your own mental model

run. Do not try to shut the mind off; instead, observe it functioning so you can detect which outputs are true versus merely familiar patterns.

13. **Journaling Practice:** Before bed each night, write down one reaction that was caught during the day and speculate on where in your training data it may originate. Consistency is key, though catching every instance will not happen due to the speed of automatic processing.

Full Transcript

The Secular Buddhism podcast is now on EightfoldPath.com. With transcripts, AI chat, and more, visit EightfoldPath.com. Hello and welcome to another episode of the Secular Buddhism podcast. I am your host, Noah Roschetta. As always, I want to remind you that you don't need to use what you learn from Buddhism to be a Buddhist. You can use what you learn to simply be a better whatever you already are. Today, I want to talk about artificial intelligence, but not quite the kind you're thinking of. I want to talk about the artificial I, the letter I as in me, myself, the one who feels like it's doing the thinking. Here's the idea I want to explore. The sense of self that feels so solid, so obviously real, works a lot like the AI systems everyone is worried about right now.

It was trained on information we never chose. It predicts, it explains, and it answers with total confidence. And most of the time, we never think to question it. So here's where we're going. First, I want to talk about why this matters, especially right now. Then I want to explain in very simple terms how AI actually learns. And you don't need to be a tech person to understand it. And once you do, I think you'll start to see the resemblance to your own mind. Then I'll explain what Buddhism means when it says the self is constructed, because that idea gets misunderstood a lot. We'll look at three things people fear AI will do to us, and notice that the self has been doing all three for as long as we've been alive, and I'll end with a practice that you can try this week.

So some of you know that a few years ago, I left teaching paragliding and started working for a data infrastructure company. That's the kind of company that builds the systems that power the internet and more and more artificial intelligence. I also went back to school to study applied artificial intelligence for a master's degree. And for the podcast's 10 year anniversary, I built a tool called NOAAI. It's trained on everything I've taught over the past decade, every episode, every book, every course. And every once in a while, I'll read one of its answers when I'm prompting it. And I'll think to myself, that sounds exactly like me. And of course, I'll realize I've never actually said that.

not in those words and sometimes not at all. But it took the way that I talk and the way that I think and produced something new that fits the pattern. And then it hit me, it isn't lying. It's doing what it was built to do. It took patterns from the past, produced the most likely next thing and set it with complete confidence. And that's where I had the slightly uncomfortable thought. Well, that's a pretty good description of the voice in my own head. So let me start with the concern or the worry because I think it's a reasonable one. A lot of people are concerned about AI right now. They're worried about jobs, about misinformation, about machines making decisions that affect our lives without anyone really understanding how those decisions got made.

And those risks are real and they deserve serious attention. I work in this field and I take them seriously. But there's another risk that we rarely talk about. It's much closer to home and I'd argue it's more dangerous than any of them. It's the mind that's going to be using the technology. So think about a megaphone for a minute. A megaphone doesn't decide what gets said. It doesn't have opinions. It takes whatever is spoken into it and makes it louder so that it can reach more people faster. If someone kind picks up the megaphone, you're going to get kindness at a high volume. If someone angry and confused picks it up, you'll get anger and confusion at a high volume. The megaphone is the same in both cases, but what changes is the person holding it. AI as a tool is a lot like that.

Whatever else it does, it amplifies the mind that's using it. A clear, caring mind with a powerful tool can do tremendous amounts of good. A confused, reactive, unexamined mind with a powerful tool produces more confusion faster. So the question that matters most isn't only what can this technology do, it's also who's holding it and do they understand their own mind? And if you look back through history, I think you'll find this pattern everywhere. Every war, every environmental disaster, every misuse of a technology started somewhere. It started in a human mind, probably a human mind that didn't see itself very clearly. Someone was absolutely sure about who they were and who the enemy was. Someone who was sure that more is always better.

The tools changed over the centuries. It's been spears, then cannons, then bombs, and now it's algorithms. But the mind holding them hasn't changed nearly as much. So AI didn't create that problem. It just inherits it. In Buddhism, there's a teaching about what are called the three poisons, greed, hatred, and delusion. Greed is wanting more and more. It's running after the things that we want. Hatred is the opposite. It's pushing away and attacking what we don't want or running away from what we don't want. And delusion is not seeing things as they really are, especially not seeing ourselves as we really are. And the idea here is that delusion is the root and greed and hatred grow out of it.

So a technology that amplifies whatever mind is using it is going to amplify all three. So that's why I think this conversation really matters. The unexamined life has always been dangerous. AI just raises the stakes. Now, here's the good news. And it's the thread for everything else in this episode. The thing holding the megaphone can be looked at. It can be understood. Because the self, like AI, is something that was put together. It's not fake. It's made. Before we talk about the self, I want to explain how these AI systems actually learn. I'm guessing a lot of people listening have used a tool like ChatGPT or at least have heard about it without maybe anyone ever explaining how tools like this actually work.

Again, you don't need any technical background for this. If you've ever sent a text message on your phone, you already understand the basic idea. So think about the keyboard on your phone. When you start typing, it suggests the next word. If you type on my, it's probably going to offer way as an option for the next possible word that you're going to type. If you type running, It offers late. How does it know? Nobody sat down and programmed your phone to know that you're always running late. It just learned that from you. Every time you typed a message, it was quietly paying attention to which words tend to follow which other words. After enough messages, it starts guessing. And it's usually right because we're all more predictable than we'd like to think.

And that's really the whole idea behind the AI tools everyone is talking about, just on an enormous scale. Instead of learning from your text messages, a system like ChatGPT learned from a huge amount of human

writing, books, articles, websites, conversations, far more than any person could read in a thousand lifetimes. And from all of that, it learned patterns. Which words tend to follow which? Which ideas tend to go together? How a question is usually answered. Then when you ask it something, it does exactly what your phone keyboard does. It predicts the most likely next word and then the next and the next until it has a whole answer. So all of that writing it learned that it learned from it has a name. It's called training data.

And here's the part I think is really important to understand. Nobody writes an AI's opinions by hand. Its opinions, its tone, its habits, its blind spots, all of that comes from the training data. If the writing it learned from leaned in a certain way, the AI leans that way too. It doesn't know that it's leaning. It just thinks that's how things are. So there are three things. I'd like you to hold on to. One, it learned from data it was given. Two, it predicts based on patterns from the past. And three, it can't tell which parts of what it learned are true and which parts are just familiar. So let's turn that around because our minds are very similar. Think about how you became you before you were old enough to evaluate anything.

You are already learning. You learned a language and with it a way of sorting the world into categories. You learned what your family valued and what they feared. You learned what got you praise and what got you in trouble. You learned what people like you were supposed to be like and what people unlike you were supposed to be like. Your culture, your religion or lack of a religion. The tone you grew up in, the things that hurt you, all of that was training data. And just like the AI, you didn't choose any of it. So imagine for a moment a kid who grows up in a house or love seems to show up when things are going well. Good grades, praise, helping out, a smile, a bad report card, silence at dinner. Nobody ever has to say, You have to earn love here. Because nobody would. It's just the kid learned that anyway, just by paying attention to what follows what. Little cues. Now fast forward 30 years, that kid is an adult who can't rest. Every time they're not being productive, they feel a low hum of anxiety that they can't explain. And if you ask them about it, they'd say, that's just who I am. I'm a hard worker.

and they'd believe it. But it's not really who they are. It's what they were trained on. The pattern was learned so early that it doesn't feel like a pattern anymore. It feels like a personality. So this is where I want to introduce a question that we'll come back to at the end of the episode. When a strong reaction shows up, the question we usually ask is, is this me or is this true? A more useful question might be, What was this trained on? Now, I want to connect this to a Buddhist teaching that often gets misunderstood. And that's the teaching of no self. When people hear that Buddhism teaches there is no self, the first reaction is usually something like, wait, so I don't exist. Then who's listening to this podcast? But that's not what it means. The teaching of no self isn't.

the claim that you aren't here or that you don't exist. You're obviously here, you have a body, you have thoughts, you have a history and a name. What the teaching points to is that the self isn't a fixed solid thing at the center of it all. It's something that gets put together moment by moment out of other things. And I think the word artificial can actually help us here. We usually hear artificial and think that means fake. But the word comes from lat and the original word meant made by skill. Art plus make. Something artificial is something that was made. Think about an artificial lake, a reservoir sitting behind a dam. It's still a real lake. You can swim in it, you can go fishing, you can drown in it. There's nothing fake about the water.

It's just a lake that was made rather than one that was always there. That's the difference. The self is the same. It's not that it's fake, it's that it's made. And the Buddha used a word that means almost exactly the same thing. The word is sankara. And it's usually translated as formations or conditioned things. Literally, it means something like put together. So the Buddha used it to describe everything that arises from causes and conditions, and the sense of self is right at the top of that list. So when Buddhism says the self is constructed, that's what it's saying. It's not that it's an illusion in the sense of not being there, but that it was put together. It was assembled from everything that came before, and it's still being assembled right now.

even as you listen to this. The idea here is that made isn't bad. Made is what makes a thing useful. An artificial lake can water a whole valley. A constructed sense of self is what lets us have a name and keep a promise

and show up for our kids and remember where we parked. The problem was never that the self is made. The problem is when we forget. that it was made. When we forget, we stop questioning it, and then it just runs. And the self isn't the only thing like this. I talked about this in the Right View episode a while back, but we live surrounded by things that humans constructed. Money, borders, our reputations, the stories about who we are and who they are. None of those existed until people made them.

And they're incredibly useful things. But when we forget that we made them, we start mistaking the story for reality, and we let the story decide how we think and act. The self is the most personal example of that. It's the constructed things that we're living inside of, and it's not fake. It's just made. So now let's go back to the worries that people have about AI. If you listen to how people talk about it, most of the fears come down to three. We're afraid AI will deceive us. We're afraid it will manipulate us. And we're afraid it will make decisions without our consent. What I want to suggest is that the self has been doing all three of those things for as long as we've been alive. And let's look at them one at a time. The first fear is that AI will deceive us.

And there's a specific way it does that, which has a name. It's called hallucination. That's when you ask an AI a question and it gives you an answer that sounds completely reasonable, confident, well written, and yet it's just flat out made up. It might invent a quote that nobody ever said or recommend a book that doesn't exist. The important thing to understand is that it isn't trying to lie. Remember, it's just predicting what a good answer would probably sound like. Sometimes what sounds right and what is right are two different things, and the AI can't tell the difference. Now it turns out our minds do this too. And there's a famous set of experiments that shows it really clearly. So our brain has two halves, a left side and a right side, and they're normally connected so that they can share information. And years ago, for some people with severe epilepsy, doctors would cut that connection as a treatment. And these people went on to live fairly normal lives, but it meant researchers could do something unusual. They could show a picture to just one half of the brain without the other half knowing about it. And for most people, the half that does the talking, the half that puts things into words is the left side. So in one experiment, They showed

the left side of a patient's brain a picture of a chicken's foot. And they showed the right side, the side that doesn't talk, a picture of a snowy yard. Then they asked him to point to matching pictures with each hand. One hand pointed to a chicken. The other hand pointed to a snow shovel. Makes sense so far.

Then they asked him, why did you pick the shovel? Now, the talking side of his brain never saw the snow. It had no idea why his hand picked a shovel. But it didn't say I don't know. Without hesitation, he said, oh, that's simple. The chicken foot goes with the chicken, and you need a shovel to clean out the chicken shed. Now, he wasn't lying. He believed it. His brain made up a reasonable sounding explanation on the spot and handed it to him as the truth. That is a hallucination. And the researchers came to see that this isn't something that only happens in unusual brains. It's something the explaining part of the mind does all the time in all of us. We do something and then a voice explains why we did it. It feels like the reason came first, but often it came after. And I think most of us know what this feels like when we look back honestly. We snap at someone. And a moment later, the voice has a whole case ready, right? Like they were being unreasonable. I was tired. Anyone would have reacted that way. And yes, maybe some of that is true. But the voice didn't make the decision. It's writing the explanation afterward with complete confidence. Again, the voice that explains us isn't necessarily the part that decided. So the second fear is that AI will manipulate us. It'll learn what makes us react and use that to steer what we do. Now remember how AI works. It predicts what comes next based on patterns from the past. Our minds do exactly the same thing.

and a lot of the time they use those predictions to steer us. So here's an example I think almost everyone has experienced. Let's say you get a text message from someone important to you, your boss, your partner, a close friend, and it just says, can we talk later? That's it, four words. Before you can even put the phone down, your mind has already finished the sentence for you, right? Uh-oh, something's wrong, I'm in trouble. They're upset with me. You might even start rehearsing what they're going to say and what you're going to say back. Maybe your heart rate goes up. Your stomach gets tight. All of that from four words that on their own don't say anything bad at all. That's the keyboard. Again, right? The phone keyboard. Your mind saw

the beginning of a sentence and it predicted the ending based on patterns from the past.

If you've had conversations that started that way and went badly, the prediction leans that way. And here's the manipulative part. The mind doesn't hand you that prediction labeled as prediction. It hands it to you as a fact, right? Wrapped in a feeling strong enough to change what you do next. Maybe you get defensive before the conversation even starts. Maybe you avoid that person the rest of the afternoon. In Buddhism, there's a teaching called the two arrows. The first arrow is the thing that actually happens. The second arrow is the story and the reaction that we had on top of it. In this example, the text message might not have been an arrow at all. But the second arrow was the whole thing, and our own mind is the one that fired it. And now we're experiencing all this anxiety, and we don't even know what any of this is about yet.

So there's the mind manipulating us. The third fear with AI is that AI will make decisions for us without our consent. It'll quietly run things in the background, and we won't even know that a decision was made. Well, once again, the self does this constantly. There's research suggesting that something like 40% of what we do each day isn't really decided in the moment. It's just habit. And I don't mean that in a bad way, habit is efficient. You don't want to decide from scratch how to brush your teeth every morning. But think about how often it shows up in places that matter. Somebody says something in a meeting and before you've had a chance to even consider it, you've already responded in a way that sounds a lot like how your dad used to respond to you.

Or you open your phone to check one thing and 20 minutes later you look up, not quite sure how you got there. And you didn't choose that, right? It was chosen. It was a pattern that was running. And you just went along with it. So the idea here is that a lot of our life is being run by a model we didn't train and have never really inspected. And that's the whole reason that this teaching matters. Again, it's not fake. It's made. the sense of self. Anything that was made can be looked at. So let's bring these together, because this is really the heart of what I wanted to share in today's episode. We worry rightly about a technology that could deceive us, manipulate us, and make

decisions without our consent. But we've been living with something that does all three, and most of us have never thought to question it.

We just assume the voices who we are. And then unexamined life is dangerous in at least three directions. It's dangerous to ourselves because we suffer from stories that we never question. The kid who learned that love has to be earned grows into an adult who can't rest and never once asks, where did that come from? That's a lot of suffering over a whole lifetime from one piece of training data. It's dangerous to others. A self needs something to define itself against. If I'm gonna be someone, there has to be someone that I'm not. That's where us and them comes from. And I think if you look at nearly every time one group of people has harmed another, you'll find a very confident story running underneath it about who we are and who they are.

Nobody experiences themselves as the villain. From the inside the story always makes sense. And then of course it's dangerous to the planet. A self that experiences itself as separate from everything will tend to treat everything as a resource. When I forget that I'm part of the web of life, the forests and the rivers become things out there rather than things I depend on and belong to. unexamined self behind the most powerful tools humans have ever built. That's the megaphone. All three of those dangers get louder and they travel faster. So what do we do about it? Well, here's something I find really hopeful and it comes from the AI world itself. One of the surprising things about these AI systems is that even the people who build them can't always explain why they say what they say. The patterns are buried so deep in all that training that nobody can simply open it up and read it. So there's a whole field of research devoted to exactly that problem. It's called interpretability. It's the work of looking inside an AI to understand why it produced the answer it did. And the reasoning is simple.

A system you can't see into is a system you can't fully trust. Mindfulness is interpretability turned inward. When we sit and meditate, we're not trying to become someone else. We're not trying to shut off the mind. We're trying to sit with the model that's running and watching how it works. We see a thought arise. We see a feeling. We follow it. We see the urge to react. And then slowly we start to understand our own outputs, our own patterns.

Again, I want to be clear that the goal isn't to get rid of the self. We don't want to turn the model off. It's useful. A tool like NOAAI that I built, it's a useful tool. And I use it, but I'd never let it make my decisions.

for me without checking its answers, we give that kind of caution to our tools all the time. We almost never give that kind of caution to our own minds, and we should. So here's the practice for this week, and it's built around the question from earlier. What was this trained on? So first, notice when you experience a strong reaction. Irritation. certainty, defensiveness, dread, any kind of reaction. The stronger and faster it shows up, the better. Those are the ones running on the deepest training. Next, name what it's predicting. Say to yourself, as a prediction, not as a fact, my mind is telling me that this means I'm in trouble or whatever it is the mind is telling you, right? It could be my mind is telling me they don't respect me.

Just putting my mind as telling me in front of it already creates some space. And then three, ask, what was this trained on? Where did I first learn to see it this way? Whose voice does this actually sound like? Is this true or is it just a pattern that's familiar? And that's it. You don't have to fix anything. You don't have to argue with the reaction or change the reaction. All you're doing is looking inside the model. Try it once a day. And if you want, before bed, write down one reaction that you caught. And where do you think it came from? And please don't be discouraged when you miss most of them because you will. I do. The training runs deep and it runs fast. So catching one is that's the practice working, right? Because we've got these things running.

all day long. And it's all pretty much automatic, like we're running on autopilot. So the idea here is let me at least pause it from time to time to see if I can at least detect one of these things and see where did that come from. So going back to where we started, when I was using the tool that I built, NOAAI, and it tells me something, I know that I need to check it. I know that, yes, it's fluent and confident and it's trained on the past. And that fluent and confident isn't the same as being true. And if you've used a tool like ChatGBT, you know the same. It'll give you a very confident answer, but that doesn't mean that it's right. What I realize is that I've never really given the same caution to the voice in my own head. That's the lesson I'm hoping

you'll take away from this, that the voice in your head has been making decisions for you way longer than any AI tool has.

Remember, the Buddha's last words, as they've been passed down to us, were, all conditioned things are subject to decay, so strive on with heedfulness. And that phrase, conditioned things, again, it's the same word we talked about earlier, sankara. It means the made things. So his final instruction wasn't to believe anything, it was to stay awake to realize what are the things that have been made, including the one that feels most like us, which is the self. The self isn't the enemy, and it isn't an illusion in the sense of not being real. It's not fake. It's made. And what's made can be seen, questioned, and even remade differently. So I'll leave you with a few questions. What am I sure of today that I was simply trained on to be sure of. Why am I so sure of that? Another one is, whose voice am I mistaking for my own? Do you ever recognize that that voice in your head? Maybe it's not yours. It's the voice you grew up with. It might be the voice of a parent or an authoritative figure from your youth. And then the last question, this is one I keep asking myself, what story am I living inside right now that I've never thought to question? That's all I have for today's episode.

Thank you for listening. Until next time. Whether you have a single question about today's episode or want to explore over a decade of teachings, it's all there. Visit 8foldpath.com to learn more.